

# Partitioning Delta Lake Tables

# Learning Objectives

- ▶ Partitioning definition
- ▶ How to partition Delta Lake tables
- ▶ Best practices

# Partitioning

- ▶ A partition is a subset of rows that share the same value for predefined partitioning columns

**partitioning column**  
↓

id	name	month	year
1	Adam	9	2015
2	Sarah	11	2015
3	John	5	2000
4	Ali	9	2000
5	Emma	10	2000
6	Eric	1	2023
7	Sophia	3	2023

**Partition 1**  
year=2015

**Partition 2**  
year=2020

**Partition 3**  
year=2023

# Partitioning Delta Lake Tables

**CREATE TABLE** my\_table (id INT, name STRING, year INT, month INT)

**PARTITIONED BY** (year)




**my\_table**

# Partitioning Delta Lake Tables

**CREATE TABLE** my\_table (id INT, name STRING, year INT, month INT)


**PARTITIONED BY** (year)



# Partition Skipping

```
SELECT * FROM my_table
```

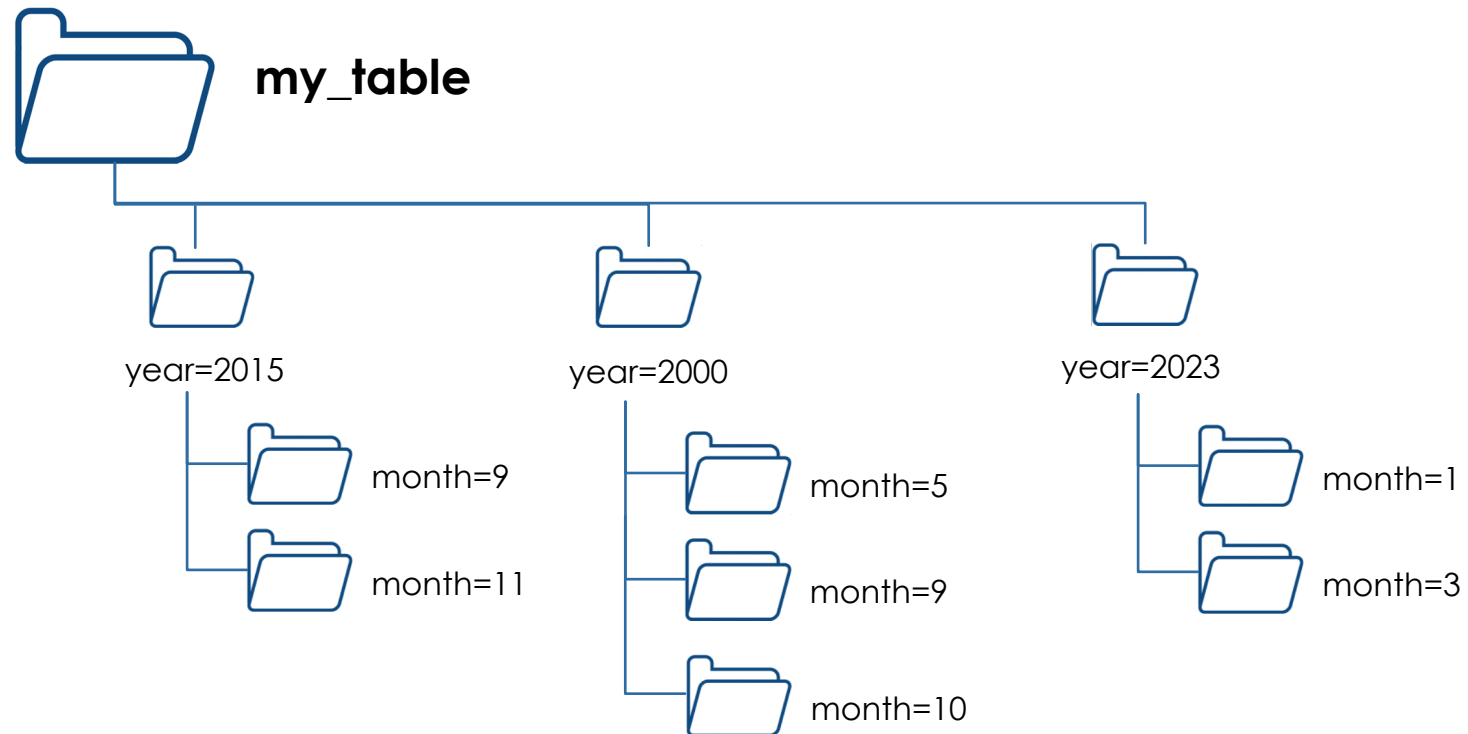
**Where year = 2023**



# Multiple Partitioning Columns

**CREATE TABLE** my\_table (id INT, name STRING, year INT, month INT)

**PARTITIONED BY** (year, month)



# Choosing partition columns

1. Total values present in a column
  - ▶ Low cardinality fields should be used for partitioning
2. Total records share a given value for a column
  - ▶ Partitions should be at least 1 GB in size
3. Records with a given value will continue to arrive indefinitely
  - ▶ Datetime fields can be used for partitioning



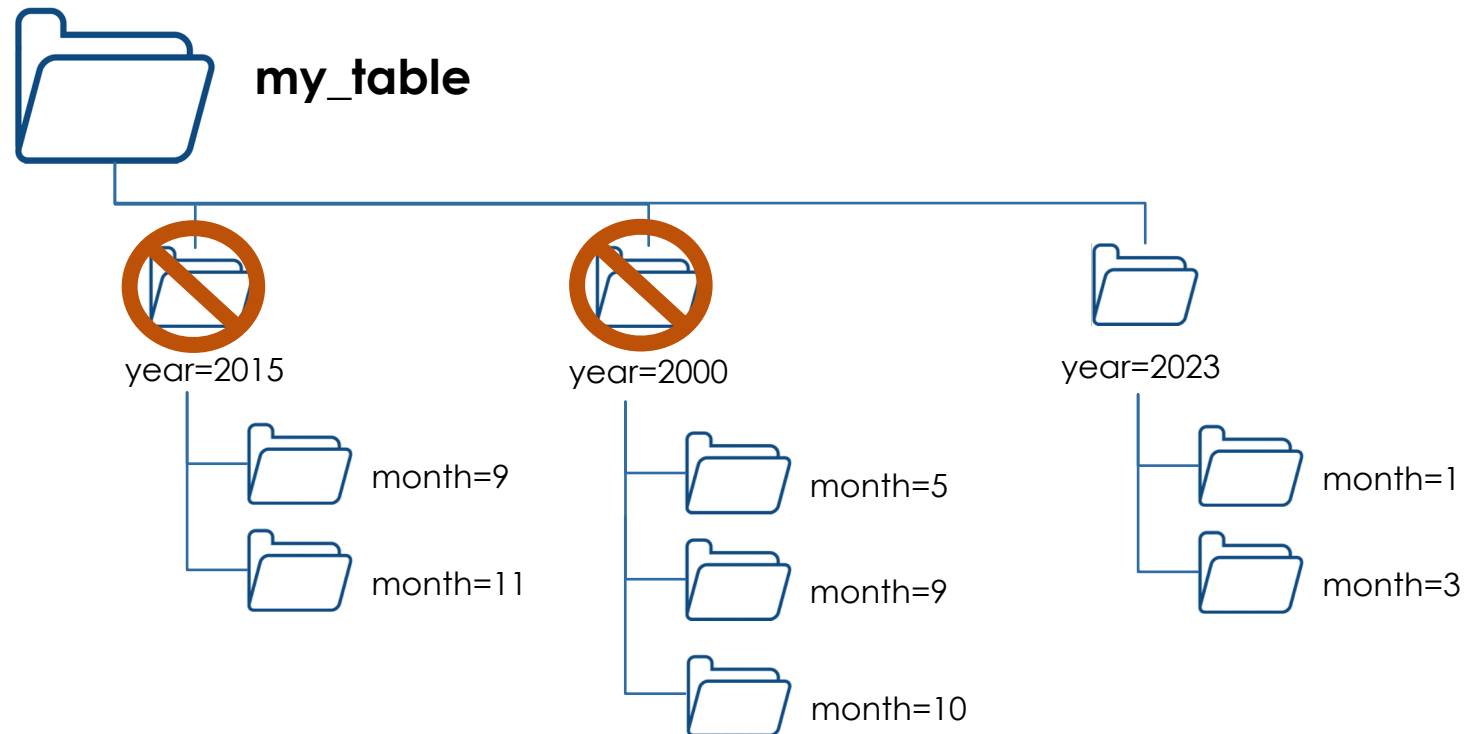
# Avoid Over-Partitioning

- ▶ Files cannot be combined or compacted across partition boundaries
  - ▶ Partitioned small tables increase storage costs and total number of files to scan.
- ▶ If most partitions < 1GB of data, the table is over-partitioned
- ▶ Data that is over-partitioned or incorrectly partitioned will suffer greatly
  - ▶ Lead to slowdowns for most general queries
  - ▶ Require a full rewrite of all data files to remedy

# Deleting at Partition Boundaries

DELETE FROM my\_table

**Where year < 2023**

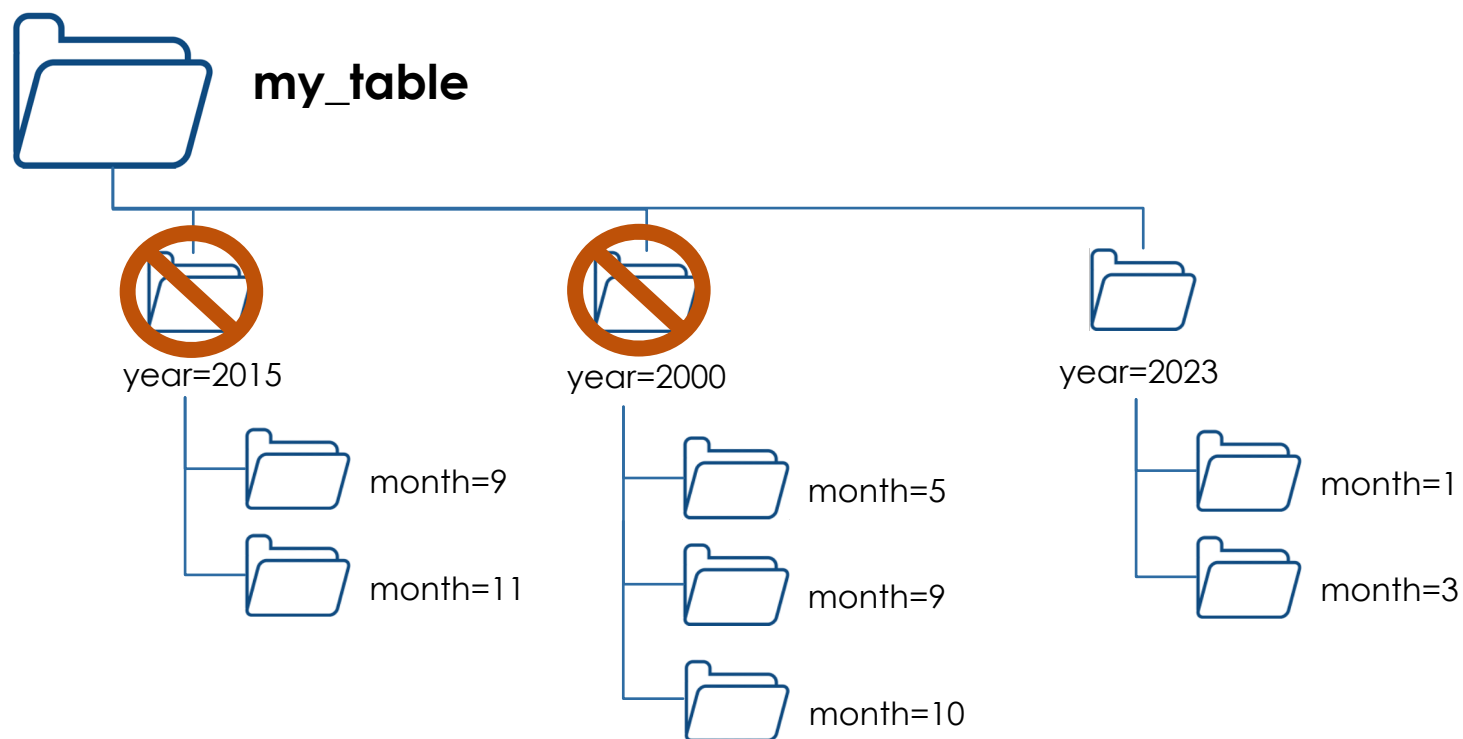


# Deleting at Partition Boundaries

```
spark.readStream
```

```
.option("ignoreDeletes", True)
```

```
.table("my_table")
```



# Deleting at Partition Boundaries

**VACUUM** my\_table

